

# Lab 2: Cats vs Dogs

In this lab, you will train a convolutional neural network to classify an image into one of two classes: "cat" or "dog". The code for the neural networks you train will be written for you, and you are not (yet!) expected to understand all provided code. However, by the end of the lab, you should be able to:

1. Understand at a high level the training loop for a machine learning model.
2. Understand the distinction between training, validation, and test data.
3. The concepts of overfitting and underfitting.
4. Investigate how different hyperparameters, such as learning rate and batch size, affect the success of training.
5. Compare an ANN (aka Multi-Layer Perceptron) with a CNN.

## What to submit

Submit a PDF file containing all your code, outputs, and write-up from parts 1-5. You can produce a PDF of your Google Colab file by going to **File > Print** and then save as PDF. The Colab instructions has more information.

**Do not submit any other files produced by your code.**

Include a link to your colab file in your submission.

Please use Google Colab to complete this assignment. If you want to use Jupyter Notebook, please complete the assignment and upload your Jupyter Notebook file to Google Colab for submission.

With Colab, you can export a PDF file using the menu option **File -> Print** and save as PDF file. **Adjust the scaling to ensure that the text is not cutoff at the margins.**

## Colab Link

Include a link to your colab file here

Colab Link: <https://colab.research.google.com/drive/15TmB7LklzWLPWP2is-iuqhFeybsKZOLZ?usp=sharing>

```
In [1]: import numpy as np
import time
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
from torch.utils.data.sampler import SubsetRandomSampler
import torchvision.transforms as transforms
```

## Part 0. Helper Functions

We will be making use of the following helper functions. You will be asked to look at and possibly modify some of these, but you are not expected to understand all of them.

You should look at the function names and read the docstrings. If you are curious, come back and explore the code *after* making some progress on the lab.

```
In [2]: #####
# Data Loading

def get_relevant_indices(dataset, classes, target_classes):
    """ Return the indices for datapoints in the dataset that belongs to the
    desired target classes, a subset of all possible classes.

    Args:
        dataset: Dataset object
        classes: A list of strings denoting the name of each class
        target_classes: A list of strings denoting the name of desired classes
                       Should be a subset of the 'classes'

    Returns:
```

```
indices: list of indices that have labels corresponding to one of the
target classes

"""
indices = []
for i in range(len(dataset)):
    # Check if the label is in the target classes
    label_index = dataset[i][1] # ex: 3
    label_class = classes[label_index] # ex: 'cat'
    if label_class in target_classes:
        indices.append(i)
return indices

def get_data_loader(target_classes, batch_size):
    """ Loads images of cats and dogs, splits the data into training, validation
and testing datasets. Returns data loaders for the three preprocessed datasets.

Args:
    target_classes: A list of strings denoting the name of the desired
                    classes. Should be a subset of the argument 'classes'
    batch_size: A int representing the number of samples per batch

Returns:
    train_loader: iterable training dataset organized according to batch size
    val_loader: iterable validation dataset organized according to batch size
    test_loader: iterable testing dataset organized according to batch size
    classes: A list of strings denoting the name of each class
"""

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
#####
# The output of torchvision datasets are PILImage images of range [0, 1].
# We transform them to Tensors of normalized range [-1, 1].
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
# Load CIFAR10 training data
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
# Get the list of indices to sample from
relevant_indices = get_relevant_indices(trainset, classes, target_classes)

# Split into train and validation
np.random.seed(1000) # Fixed numpy random seed for reproducible shuffling
np.random.shuffle(relevant_indices)
split = int(len(relevant_indices) * 0.8) #split at 80%

# split into training and validation indices
relevant_train_indices, relevant_val_indices = relevant_indices[:split], relevant_indices[split:]
train_sampler = SubsetRandomSampler(relevant_train_indices)
train_loader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           num_workers=1, sampler=train_sampler)
val_sampler = SubsetRandomSampler(relevant_val_indices)
val_loader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                          num_workers=1, sampler=val_sampler)

# Load CIFAR10 testing data
testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                        download=True, transform=transform)
# Get the list of indices to sample from
relevant_test_indices = get_relevant_indices(testset, classes, target_classes)
test_sampler = SubsetRandomSampler(relevant_test_indices)
test_loader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                           num_workers=1, sampler=test_sampler)

return train_loader, val_loader, test_loader, classes

#####
# Training
def get_model_name(name, batch_size, learning_rate, epoch):
    """ Generate a name for the model consisting of all the hyperparameter values

Args:
    config: Configuration object containing the hyperparameters
Returns:
    path: A string with the hyperparameter name and value concatenated
"""
path = "model_{0}_bs{1}_lr{2}_epoch{3}".format(name,
```

```

learning_rate,
epoch)

return path

def normalize_label(labels):
    """
    Given a tensor containing 2 possible values, normalize this to 0/1

    Args:
        labels: a 1D tensor containing two possible scalar values
    Returns:
        A tensor normalize to 0/1 value
    """
    max_val = torch.max(labels)
    min_val = torch.min(labels)
    norm_labels = (labels - min_val)/(max_val - min_val)
    return norm_labels

def evaluate(net, loader, criterion):
    """ Evaluate the network on the validation set.

    Args:
        net: PyTorch neural network object
        loader: PyTorch data loader for the validation set
        criterion: The loss function
    Returns:
        err: A scalar for the avg classification error over the validation set
        loss: A scalar for the average loss function over the validation set
    """
    total_loss = 0.0
    total_err = 0.0
    total_epoch = 0
    for i, data in enumerate(loader, 0):
        inputs, labels = data
        labels = normalize_label(labels) # Convert labels to 0/1
        outputs = net(inputs)
        loss = criterion(outputs, labels.float())
        corr = (outputs > 0.0).squeeze().long() != labels
        total_err += int(corr.sum())
        total_loss += loss.item()
        total_epoch += len(labels)
    err = float(total_err) / total_epoch
    loss = float(total_loss) / (i + 1)
    return err, loss

#####
# Training Curve
def plot_training_curve(path):
    """ Plots the training curve for a model run, given the csv files
    containing the train/validation error/loss.

    Args:
        path: The base path of the csv files produced during training
    """
    import matplotlib.pyplot as plt
    train_err = np.loadtxt("{}_train_err.csv".format(path))
    val_err = np.loadtxt("{}_val_err.csv".format(path))
    train_loss = np.loadtxt("{}_train_loss.csv".format(path))
    val_loss = np.loadtxt("{}_val_loss.csv".format(path))
    plt.title("Train vs Validation Error")
    n = len(train_err) # number of epochs
    plt.plot(range(1,n+1), train_err, label="Train")
    plt.plot(range(1,n+1), val_err, label="Validation")
    plt.xlabel("Epoch")
    plt.ylabel("Error")
    plt.legend(loc='best')
    plt.show()
    plt.title("Train vs Validation Loss")
    plt.plot(range(1,n+1), train_loss, label="Train")
    plt.plot(range(1,n+1), val_loss, label="Validation")
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.legend(loc='best')
    plt.show()

```

## Part 1. Visualizing the Data [7 pt]

We will make use of some of the CIFAR-10 data set, which consists of colour images of size 32x32 pixels belonging to 10 categories. You can find out more about the dataset at <https://www.cs.toronto.edu/~kriz/cifar.html>

For this assignment, we will only be using the cat and dog categories. We have included code that automatically downloads the dataset the first time that the main script is run.

```
In [3]: # This will download the CIFAR-10 dataset to a folder called "data"
# the first time you run this code.
train_loader, val_loader, test_loader, classes = get_data_loader(
    target_classes=["cat", "dog"],
    batch_size=1) # One image per batch
```

100%|██████████| 170M/170M [00:01<00:00, 105MB/s]

## Part (a) -- 1 pt


Visualize some of the data by running the code below. Include the visualization in your writeup.

(You don't need to submit anything else.)

```
In [4]: import matplotlib.pyplot as plt

k = 0
for images, labels in train_loader:
    # since batch_size = 1, there is only 1 image in `images`
    image = images[0]
    # place the colour channel at the end, instead of at the beginning
    img = np.transpose(image, [1,2,0])
    # normalize pixel intensity values to [0, 1]
    img = img / 2 + 0.5
    plt.subplot(3, 5, k+1)
    plt.axis('off')
    plt.imshow(img)

    k += 1
    if k > 14:
        break
```



## Part (b) -- 3 pt

How many training examples do we have for the combined `cat` and `dog` classes? What about validation examples? What about test examples?

```
In [5]: print("Number of training examples: {}".format(len(train_loader)))
print("Number of validation examples: {}".format(len(val_loader)))
print("Number of test examples: {}".format(len(test_loader)))
```

Number of training examples: 8000  
Number of validation examples: 2000  
Number of test examples: 2000

## Part (c) -- 3pt

Why do we need a validation set when training our model? What happens if we judge the performance of our models using the training set loss/error instead of the validation set loss/error?

**Answer:**

We need a validation set to get an unbiased evaluation of the model during training and to detect overfitting/underfitting. Using the training set error alone can be misleading as the model will naturally perform well on data it has already seen, potentially leading to overfitting where the model doesn't generalize well to new data. The validation set acts as a proxy for unseen data.

## Part 2. Training [15 pt]

We define two neural networks, a `LargeNet` and `SmallNet`. We'll be training the networks in this section.

You won't understand fully what these networks are doing until the next few classes, and that's okay. For this assignment, please focus on learning how to train networks, and how hyperparameters affect training.

```
In [6]: class LargeNet(nn.Module):
        def __init__(self):
            super(LargeNet, self).__init__()
            self.name = "large"
            self.conv1 = nn.Conv2d(3, 5, 5)
            self.pool = nn.MaxPool2d(2, 2)
            self.conv2 = nn.Conv2d(5, 10, 5)
            self.fc1 = nn.Linear(10 * 5 * 5, 32)
            self.fc2 = nn.Linear(32, 1)

        def forward(self, x):
            x = self.pool(F.relu(self.conv1(x)))
            x = self.pool(F.relu(self.conv2(x)))
            x = x.view(-1, 10 * 5 * 5)
            x = F.relu(self.fc1(x))
            x = self.fc2(x)
            x = x.squeeze(1) # Flatten to [batch_size]
            return x
```

```
In [7]: class SmallNet(nn.Module):
        def __init__(self):
            super(SmallNet, self).__init__()
            self.name = "small"
            self.conv = nn.Conv2d(3, 5, 3)
            self.pool = nn.MaxPool2d(2, 2)
            self.fc = nn.Linear(5 * 7 * 7, 1)

        def forward(self, x):
            x = self.pool(F.relu(self.conv(x)))
            x = self.pool(x)
            x = x.view(-1, 5 * 7 * 7)
            x = self.fc(x)
            x = x.squeeze(1) # Flatten to [batch_size]
            return x
```

```
In [8]: small_net = SmallNet()
        large_net = LargeNet()
```

## Part (a) -- 2pt

The methods `small_net.parameters()` and `large_net.parameters()` produces an iterator of all the trainable parameters of the network. These parameters are torch tensors containing many scalar values.

We haven't learned how the parameters in these high-dimensional tensors will be used, but we should be able to count the number of parameters. Measuring the number of parameters in a network is one way of measuring the "size" of a network.

What is the total number of parameters in `small_net` and in `large_net`? (Hint: how many numbers are in each tensor?)

```
In [9]: print("small_net:")
        for param in small_net.parameters():
            print(param.shape)
```

```
print()

print("large_net")
for param in large_net.parameters():
    print(param.shape)
```

```
small_net:
torch.Size([5, 3, 3, 3])
torch.Size([5])
torch.Size([1, 245])
torch.Size([1])
```

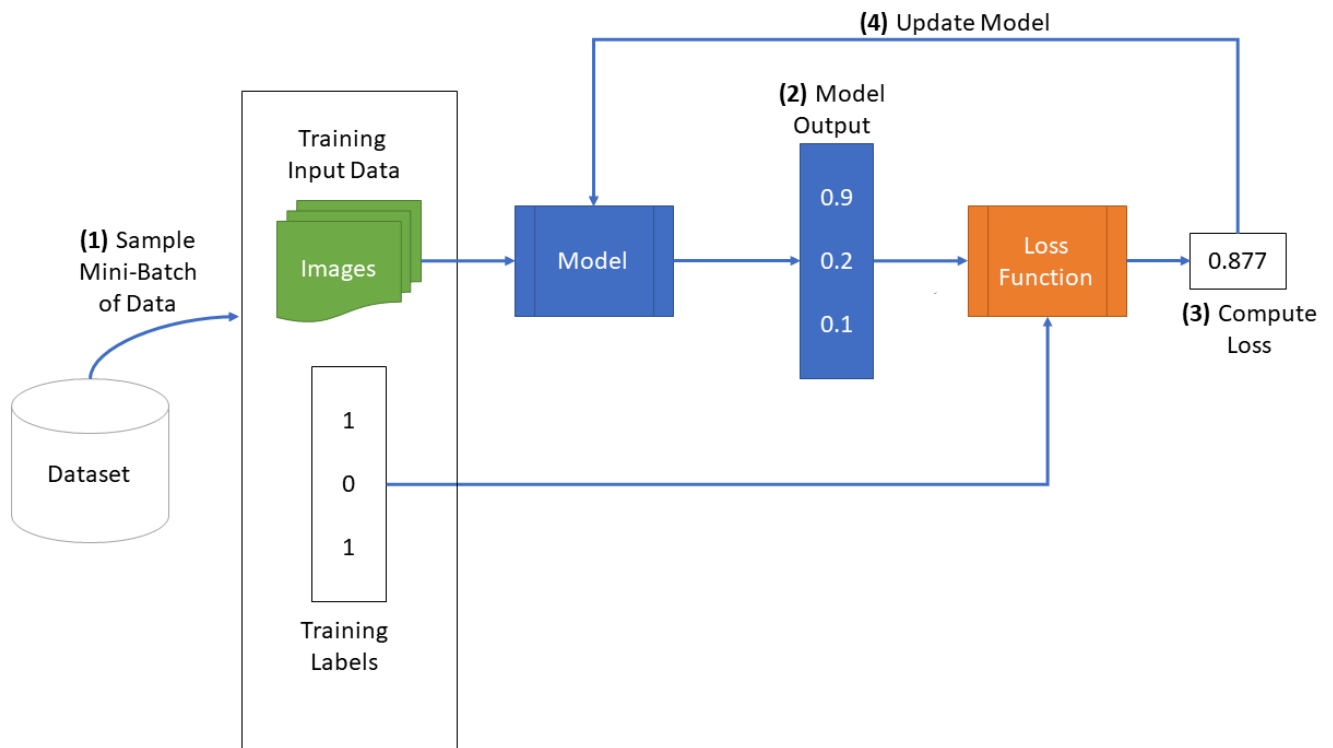
```
large_net
torch.Size([5, 3, 5, 5])
torch.Size([5])
torch.Size([10, 5, 5, 5])
torch.Size([10])
torch.Size([32, 250])
torch.Size([32])
torch.Size([1, 32])
torch.Size([1])
```

**Answer:**

- $(5*3*3*3)+(5)+(1*245)+(1)=386$
- $(5*3*5*5)+(5)+(10*5*5*5)+(10)+(32*250)+(32)+(1*32)+(1)=9705$
- There are **386** parameters in the small net, and **9705** parameters in the large net.

## The function train\_net

The function `train_net` below takes an untrained neural network (like `small_net` and `large_net`) and several other parameters. You should be able to understand how this function works. The figure below shows the high level training loop for a machine learning model:



```
In [10]: def train_net(net, batch_size=64, learning_rate=0.01, num_epochs=30):
#####
# Train a classifier on cats vs dogs
target_classes = ["cat", "dog"]
#####
# Fixed PyTorch random seed for reproducible result
torch.manual_seed(1000)
#####
```

```

# Obtain the PyTorch data loader objects to load batches of the datasets
train_loader, val_loader, test_loader, classes = get_data_loader(
    target_classes, batch_size)
#####
# Define the Loss function and optimizer
# The loss function will be Binary Cross Entropy (BCE). In this case we
# will use the BCEWithLogitsLoss which takes unnormalized output from
# the neural network and scalar label.
# Optimizer will be SGD with Momentum.
criterion = nn.BCEWithLogitsLoss()
optimizer = optim.SGD(net.parameters(), lr=learning_rate, momentum=0.9)
#####
# Set up some numpy arrays to store the training/test loss/erruracy
train_err = np.zeros(num_epochs)
train_loss = np.zeros(num_epochs)
val_err = np.zeros(num_epochs)
val_loss = np.zeros(num_epochs)
#####
# Train the network
# Loop over the data iterator and sample a new batch of training data
# Get the output from the network, and optimize our loss function.
start_time = time.time()
for epoch in range(num_epochs): # loop over the dataset multiple times
    total_train_loss = 0.0
    total_train_err = 0.0
    total_epoch = 0
    for i, data in enumerate(train_loader, 0):
        # Get the inputs
        inputs, labels = data
        labels = normalize_label(labels) # Convert labels to 0/1
        # Zero the parameter gradients
        optimizer.zero_grad()
        # Forward pass, backward pass, and optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels.float())
        loss.backward()
        optimizer.step()
        # Calculate the statistics
        corr = (outputs > 0.0).squeeze().long() != labels
        total_train_err += int(corr.sum())
        total_train_loss += loss.item()
        total_epoch += len(labels)
    train_err[epoch] = float(total_train_err) / total_epoch
    train_loss[epoch] = float(total_train_loss) / (i+1)
    val_err[epoch], val_loss[epoch] = evaluate(net, val_loader, criterion)
    print(("Epoch {}: Train err: {}, Train loss: {} |"+
        "Validation err: {}, Validation loss: {}".format(
            epoch + 1,
            train_err[epoch],
            train_loss[epoch],
            val_err[epoch],
            val_loss[epoch]))
    # Save the current model (checkpoint) to a file
    model_path = get_model_name(net.name, batch_size, learning_rate, epoch)
    torch.save(net.state_dict(), model_path)
print('Finished Training')
end_time = time.time()
elapsed_time = end_time - start_time
print("Total time elapsed: {:.2f} seconds".format(elapsed_time))
# Write the train/test loss/err into CSV file for plotting later
epochs = np.arange(1, num_epochs + 1)
np.savetxt("{}_train_err.csv".format(model_path), train_err)
np.savetxt("{}_train_loss.csv".format(model_path), train_loss)
np.savetxt("{}_val_err.csv".format(model_path), val_err)
np.savetxt("{}_val_loss.csv".format(model_path), val_loss)

```

## Part (b) -- 1pt

The parameters to the function `train_net` are hyperparameters of our neural network. We made these hyperparameters easy to modify so that we can tune them later on.

What are the default values of the parameters `batch_size`, `learning_rate`, and `num_epochs` ?

```

In [11]: batch_size=64
         learning_rate=0.01

```

```
num_epochs=30
```

## Part (c) -- 3 pt

What files are written to disk when we call `train_net` with `small_net`, and train for 5 epochs? Provide a list of all the files written to disk, and what information the files contain.

```
In [12]: train_net(small_net,num_epochs=5)
```

```
Epoch 1: Train err: 0.432375, Train loss: 0.6764058194160462 |Validation err: 0.3785, Validation loss: 0.6579908803
105354
Epoch 2: Train err: 0.370625, Train loss: 0.6472312378883361 |Validation err: 0.384, Validation loss: 0.65966347977
51904
Epoch 3: Train err: 0.355625, Train loss: 0.6366147956848145 |Validation err: 0.353, Validation loss: 0.62876553088
42659
Epoch 4: Train err: 0.341875, Train loss: 0.6221265025138855 |Validation err: 0.3495, Validation loss: 0.6207225974
649191
Epoch 5: Train err: 0.332125, Train loss: 0.6116174993515014 |Validation err: 0.333, Validation loss: 0.61577494069
93389
Finished Training
Total time elapsed: 23.77 seconds
```

**Answer:**

- checkpoints:
  - `model_small_bs64_lr0.01_epoch0`: checkpoint at epoch 0
  - `model_small_bs64_lr0.01_epoch1`: checkpoint at epoch 1
  - `model_small_bs64_lr0.01_epoch2`: checkpoint at epoch 2
  - `model_small_bs64_lr0.01_epoch3`: checkpoint at epoch 3
  - `model_small_bs64_lr0.01_epoch4`: checkpoint at epoch 4
- Train error/loss:
  - `model_small_bs64_lr0.01_epoch4_train_err.csv`: training error
  - `model_small_bs64_lr0.01_epoch4_train_loss.csv`: training loss
  - `model_small_bs64_lr0.01_epoch4_val_err.csv`: validation error
  - `model_small_bs64_lr0.01_epoch4_val_loss.csv`: validation loss

## Part (d) -- 2pt

Train both `small_net` and `large_net` using the function `train_net` and its default parameters. The function will write many files to disk, including a model checkpoint (saved values of model weights) at the end of each epoch.

If you are using Google Colab, you will need to mount Google Drive so that the files generated by `train_net` gets saved. We will be using these files in part (d). (See the Google Colab tutorial for more information about this.)

Report the total time elapsed when training each network. Which network took longer to train? Why?

```
In [14]: # Since the function writes files to disk, you will need to mount
# your Google Drive. If you are working on the lab locally, you
# can comment out this code.

from google.colab import drive
drive.mount('/content/gdrive')
```

Mounted at `/content/gdrive`

```
In [15]: print("Training small_net")
small_net = SmallNet()
train_net(small_net)

print("Training large_net")
large_net = LargeNet()
train_net(large_net)
```



Training small\_net  
Epoch 1: Train err: 0.446375, Train loss: 0.6813716740608216 |Validation err: 0.3865, Validation loss: 0.660299738869071  
Epoch 2: Train err: 0.37325, Train loss: 0.6497629313468933 |Validation err: 0.3845, Validation loss: 0.6575995795428753  
Epoch 3: Train err: 0.360125, Train loss: 0.6389007897377014 |Validation err: 0.3495, Validation loss: 0.6291371881961823  
Epoch 4: Train err: 0.346375, Train loss: 0.6246497564315796 |Validation err: 0.356, Validation loss: 0.6221635770052671  
Epoch 5: Train err: 0.334125, Train loss: 0.6154210748672485 |Validation err: 0.328, Validation loss: 0.618933217599988  
Epoch 6: Train err: 0.31775, Train loss: 0.6036666023731232 |Validation err: 0.3385, Validation loss: 0.6093297032639384  
Epoch 7: Train err: 0.315875, Train loss: 0.5944248065948486 |Validation err: 0.328, Validation loss: 0.5973024442791939  
Epoch 8: Train err: 0.308125, Train loss: 0.5829446413516999 |Validation err: 0.3085, Validation loss: 0.5884840954095125  
Epoch 9: Train err: 0.302875, Train loss: 0.5804586551189422 |Validation err: 0.312, Validation loss: 0.5844892263412476  
Epoch 10: Train err: 0.29825, Train loss: 0.5729375307559967 |Validation err: 0.3095, Validation loss: 0.5786221697926521  
Epoch 11: Train err: 0.28825, Train loss: 0.5630653114318848 |Validation err: 0.313, Validation loss: 0.5817071627825499  
Epoch 12: Train err: 0.292875, Train loss: 0.5564581718444824 |Validation err: 0.312, Validation loss: 0.5859681870788336  
Epoch 13: Train err: 0.288625, Train loss: 0.5559324417114258 |Validation err: 0.3045, Validation loss: 0.5765548050403595  
Epoch 14: Train err: 0.28, Train loss: 0.5471521301269531 |Validation err: 0.3115, Validation loss: 0.5721538551151752  
Epoch 15: Train err: 0.285125, Train loss: 0.5481183273792267 |Validation err: 0.3045, Validation loss: 0.5623642448335886  
Epoch 16: Train err: 0.29175, Train loss: 0.5540513117313385 |Validation err: 0.31, Validation loss: 0.5772618055343628  
Epoch 17: Train err: 0.28225, Train loss: 0.5475915327072144 |Validation err: 0.3, Validation loss: 0.5682430267333984  
Epoch 18: Train err: 0.2805, Train loss: 0.5438903319835663 |Validation err: 0.3185, Validation loss: 0.577332322485745  
Epoch 19: Train err: 0.275625, Train loss: 0.5399831240177154 |Validation err: 0.33, Validation loss: 0.6077376063913107  
Epoch 20: Train err: 0.272875, Train loss: 0.5384533171653747 |Validation err: 0.296, Validation loss: 0.5776732238009572  
Epoch 21: Train err: 0.27675, Train loss: 0.5400643782615662 |Validation err: 0.302, Validation loss: 0.5675954343751073  
Epoch 22: Train err: 0.277375, Train loss: 0.5396291089057922 |Validation err: 0.288, Validation loss: 0.5697800805792212  
Epoch 23: Train err: 0.2735, Train loss: 0.5354287037849427 |Validation err: 0.302, Validation loss: 0.567152694799006  
Epoch 24: Train err: 0.272625, Train loss: 0.5356638813018799 |Validation err: 0.2995, Validation loss: 0.5875891577452421  
Epoch 25: Train err: 0.272875, Train loss: 0.534599012374878 |Validation err: 0.2965, Validation loss: 0.563906635157764  
Epoch 26: Train err: 0.269875, Train loss: 0.5316087663173675 |Validation err: 0.2965, Validation loss: 0.56999272108078  
Epoch 27: Train err: 0.27, Train loss: 0.5297632284164429 |Validation err: 0.301, Validation loss: 0.5788038447499275  
Epoch 28: Train err: 0.270625, Train loss: 0.535400269985199 |Validation err: 0.2995, Validation loss: 0.565693317912519  
Epoch 29: Train err: 0.272375, Train loss: 0.531760558128357 |Validation err: 0.295, Validation loss: 0.5849149720743299  
Epoch 30: Train err: 0.271, Train loss: 0.537471985578537 |Validation err: 0.315, Validation loss: 0.5815007202327251  
Finished Training  
Total time elapsed: 141.55 seconds  
Training large\_net  
Epoch 1: Train err: 0.444625, Train loss: 0.6900211324691773 |Validation err: 0.4305, Validation loss: 0.6807530857622623  
Epoch 2: Train err: 0.418875, Train loss: 0.6781997265815735 |Validation err: 0.4125, Validation loss: 0.6741086076945066  
Epoch 3: Train err: 0.39825, Train loss: 0.6658710074424744 |Validation err: 0.394, Validation loss: 0.6518888883292675  
Epoch 4: Train err: 0.3745, Train loss: 0.6488627505302429 |Validation err: 0.408, Validation loss: 0.6639942992478609  
Epoch 5: Train err: 0.35475, Train loss: 0.633074896812439 |Validation err: 0.352, Validation loss: 0.6291442140936852  
Epoch 6: Train err: 0.340625, Train loss: 0.6164079172611237 |Validation err: 0.3425, Validation loss: 0.6155355125665665  
Epoch 7: Train err: 0.3265, Train loss: 0.6007295575141907 |Validation err: 0.3395, Validation loss: 0.609989359974

```

8611
Epoch 8: Train err: 0.313375, Train loss: 0.5828288238048553 |Validation err: 0.327, Validation loss: 0.59817310236
3944
Epoch 9: Train err: 0.3095, Train loss: 0.5759536104202271 |Validation err: 0.3205, Validation loss: 0.595438154414
2962
Epoch 10: Train err: 0.29475, Train loss: 0.5617487483024597 |Validation err: 0.3145, Validation loss: 0.5929642692
20829
Epoch 11: Train err: 0.281, Train loss: 0.5484012789726257 |Validation err: 0.3155, Validation loss: 0.608386868610
9781
Epoch 12: Train err: 0.275875, Train loss: 0.5395474998950959 |Validation err: 0.3175, Validation loss: 0.597383063
2865429
Epoch 13: Train err: 0.271875, Train loss: 0.5274702501296997 |Validation err: 0.298, Validation loss: 0.5836638603
359461
Epoch 14: Train err: 0.260875, Train loss: 0.5171499395370484 |Validation err: 0.297, Validation loss: 0.5946236466
988921
Epoch 15: Train err: 0.25775, Train loss: 0.5102220706939697 |Validation err: 0.2995, Validation loss: 0.5893601188
43615
Epoch 16: Train err: 0.24725, Train loss: 0.5006417303085328 |Validation err: 0.299, Validation loss: 0.59780285693
70508
Epoch 17: Train err: 0.245375, Train loss: 0.4954251596927643 |Validation err: 0.3045, Validation loss: 0.586751978
8444042
Epoch 18: Train err: 0.2335, Train loss: 0.48036268329620363 |Validation err: 0.304, Validation loss: 0.61384835559
87477
Epoch 19: Train err: 0.233, Train loss: 0.47708055830001833 |Validation err: 0.333, Validation loss: 0.624375478364
5272
Epoch 20: Train err: 0.22875, Train loss: 0.46485390305519103 |Validation err: 0.3115, Validation loss: 0.603004706
0921788
Epoch 21: Train err: 0.22, Train loss: 0.45495913982391356 |Validation err: 0.2905, Validation loss: 0.602344034239
6498
Epoch 22: Train err: 0.217, Train loss: 0.4491056790351868 |Validation err: 0.2965, Validation loss: 0.618036177940
6667
Epoch 23: Train err: 0.217125, Train loss: 0.44135768175125123 |Validation err: 0.301, Validation loss: 0.604913904
3316245
Epoch 24: Train err: 0.203125, Train loss: 0.4262930202484131 |Validation err: 0.3165, Validation loss: 0.667187351
7334461
Epoch 25: Train err: 0.198375, Train loss: 0.4121157767726746 |Validation err: 0.312, Validation loss: 0.634801496
7516065
Epoch 26: Train err: 0.1875, Train loss: 0.39792571806907656 |Validation err: 0.295, Validation loss: 0.64646175038
06949
Epoch 27: Train err: 0.184375, Train loss: 0.3930838825702667 |Validation err: 0.31, Validation loss: 0.64248665794
73019
Epoch 28: Train err: 0.176, Train loss: 0.38016000711917874 |Validation err: 0.3105, Validation loss: 0.65657526999
71199
Epoch 29: Train err: 0.166625, Train loss: 0.3643355256319046 |Validation err: 0.33, Validation loss: 0.78958085179
32892
Epoch 30: Train err: 0.158375, Train loss: 0.3545009701251984 |Validation err: 0.303, Validation loss: 0.6914397869
259119
Finished Training
Total time elapsed: 164.74 seconds

```

### Answer:

Time for training small network is 141.55s, and for large network is 164.74s. Large network has much more parameters to update so it takes longer.

## Part (e) - 2pt

Use the function `plot_training_curve` to display the trajectory of the training/validation error and the training/validation loss. You will need to use the function `get_model_name` to generate the argument to the `plot_training_curve` function.

Do this for both the small network and the large network. Include both plots in your writeup.

```


In [16]: #model_path = get_model_name("small", batch_size=??, learning_rate=??, epoch=29)
small_path = get_model_name("small", batch_size=64, learning_rate=0.01, epoch=29)
large_path = get_model_name("large", batch_size=64, learning_rate=0.01, epoch=29)

print("Small Network Plot: ")
plot_training_curve(small_path)
print("Large Network Plot: ")
plot_training_curve(large_path)


```

Small Network Plot:



Train vs Validation Error



Train vs Validation Loss



Large Network Plot:



### Part (f) - 5pt

Describe what you notice about the training curve. How do the curves differ for `small_net` and `large_net` ? Identify any occurrences of underfitting and overfitting.

**Answer:**

At first, both networks' training and validation loss and error decrease steadily as the optimizer refines parameters, but around epoch 5 for the large net and epoch 10 for the small net the validation curves stop improving. The small network's validation error and loss then plateau around the same level—indicating limited capacity and mild underfitting—while the large network's validation loss actually rises even as its training error falls, signaling clear overfitting (i.e., the model memorizes the training set and fails to generalize).

## Part 3. Optimization Parameters [12 pt]


For this section, we will work with `large_net` only.

### Part (a) - 3pt

Train `large_net` with all default parameters, except set `learning_rate=0.001`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *lowering* the learning rate.

```
In [17]: # Note: When we re-construct the model, we start the training
# with *random weights*. If we omit this code, the values of
# the weights will still be the previously trained values.
large_net = LargeNet()
train_net(large_net, learning_rate=0.001)
large_path = get_model_name("large", batch_size = 64, learning_rate = 0.001, epoch = 29)
plot_training_curve(large_path)
```

```
Epoch 1: Train err: 0.47625, Train loss: 0.6928360004425049 |Validation err: 0.467, Validation loss: 0.692468658089
6378
Epoch 2: Train err: 0.448625, Train loss: 0.6922589716911316 |Validation err: 0.4305, Validation loss: 0.6916493363
678455
Epoch 3: Train err: 0.43575, Train loss: 0.6916067404747009 |Validation err: 0.4285, Validation loss: 0.69085445255
04112
Epoch 4: Train err: 0.430125, Train loss: 0.6908613877296448 |Validation err: 0.424, Validation loss: 0.68965969607
234
Epoch 5: Train err: 0.434125, Train loss: 0.6899198365211486 |Validation err: 0.4195, Validation loss: 0.6886942796
40913
Epoch 6: Train err: 0.435875, Train loss: 0.6887419748306275 |Validation err: 0.4195, Validation loss: 0.6867837496
101856
Epoch 7: Train err: 0.436625, Train loss: 0.6873781814575195 |Validation err: 0.4185, Validation loss: 0.6851996649
056673
Epoch 8: Train err: 0.43725, Train loss: 0.6859267811775207 |Validation err: 0.4115, Validation loss: 0.68319912441
07485
Epoch 9: Train err: 0.424375, Train loss: 0.6844044809341431 |Validation err: 0.411, Validation loss: 0.68088660389
18495
Epoch 10: Train err: 0.42425, Train loss: 0.6828490204811096 |Validation err: 0.408, Validation loss: 0.67835014313
4594
Epoch 11: Train err: 0.425375, Train loss: 0.6812362918853759 |Validation err: 0.4125, Validation loss: 0.678020710
1255655
Epoch 12: Train err: 0.420125, Train loss: 0.6796324462890625 |Validation err: 0.4125, Validation loss: 0.675314700
2309561
Epoch 13: Train err: 0.414875, Train loss: 0.6777921686172486 |Validation err: 0.415, Validation loss: 0.6757054049
521685
Epoch 14: Train err: 0.41225, Train loss: 0.6761114087104797 |Validation err: 0.412, Validation loss: 0.67397416010
499
Epoch 15: Train err: 0.409, Train loss: 0.6744724254608154 |Validation err: 0.415, Validation loss: 0.6706814523786
306
Epoch 16: Train err: 0.406375, Train loss: 0.6727445869445801 |Validation err: 0.4105, Validation loss: 0.670771589
5026922
Epoch 17: Train err: 0.401375, Train loss: 0.67130890417099 |Validation err: 0.404, Validation loss: 0.667154807597
3988
Epoch 18: Train err: 0.399375, Train loss: 0.6696756863594056 |Validation err: 0.4055, Validation loss: 0.664682278
4096003
Epoch 19: Train err: 0.40075, Train loss: 0.6679100017547608 |Validation err: 0.396, Validation loss: 0.66550736874
34196
Epoch 20: Train err: 0.39225, Train loss: 0.665790048122406 |Validation err: 0.4045, Validation loss: 0.66260510496
79518
Epoch 21: Train err: 0.389625, Train loss: 0.6646309685707092 |Validation err: 0.3945, Validation loss: 0.660683380
4398775
Epoch 22: Train err: 0.38875, Train loss: 0.6623744034767151 |Validation err: 0.393, Validation loss: 0.66170048713
68408
Epoch 23: Train err: 0.384125, Train loss: 0.6601551241874695 |Validation err: 0.3975, Validation loss: 0.657400136
8135214
Epoch 24: Train err: 0.382375, Train loss: 0.6584072341918945 |Validation err: 0.3865, Validation loss: 0.656140377
7450323
Epoch 25: Train err: 0.378875, Train loss: 0.6555052490234375 |Validation err: 0.3885, Validation loss: 0.655285593
1222439
Epoch 26: Train err: 0.37675, Train loss: 0.6531328277587891 |Validation err: 0.387, Validation loss: 0.65318292379
37927
Epoch 27: Train err: 0.375125, Train loss: 0.6503917617797852 |Validation err: 0.387, Validation loss: 0.6519917994
737625
Epoch 28: Train err: 0.371375, Train loss: 0.647677414894104 |Validation err: 0.3875, Validation loss: 0.6483855955
302715
Epoch 29: Train err: 0.368, Train loss: 0.6451585369110108 |Validation err: 0.382, Validation loss: 0.6459537353366
613
Epoch 30: Train err: 0.362875, Train loss: 0.6423822708129883 |Validation err: 0.3785, Validation loss: 0.643928943
2018995
Finished Training
Total time elapsed: 166.52 seconds
```



**Answer:**



It runs 1.78 seconds longer than the default setting, this overall runtime are very similar. With a lower learning rate, training error and loss both increase, and validation error and loss rise as well. However, the training and validation curves are closer together, indicating that overfitting, as observed when the learning rate is 0.01, is avoided.

### Part (b) - 3pt

Train `large_net` with all default parameters, except set `learning_rate=0.1`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *increasing* the learning rate.

```
In [19]: large_net = LargeNet()
train_net(large_net, learning_rate=0.1)
large_path = get_model_name("large", batch_size=64, learning_rate=0.1, epoch=29)
plot_training_curve(large_path)
```

Epoch 1: Train err: 0.468, Train loss: 0.6859708547592163 |Validation err: 0.3915, Validation loss: 0.6671512890607119  
Epoch 2: Train err: 0.38175, Train loss: 0.6535712971687316 |Validation err: 0.416, Validation loss: 0.6825554035604  
Epoch 3: Train err: 0.355625, Train loss: 0.6300361003875733 |Validation err: 0.3315, Validation loss: 0.6121374126523733  
Epoch 4: Train err: 0.33875, Train loss: 0.6164356069564819 |Validation err: 0.329, Validation loss: 0.6106978859752417  
Epoch 5: Train err: 0.331625, Train loss: 0.6028243894577027 |Validation err: 0.322, Validation loss: 0.602822014130652  
Epoch 6: Train err: 0.302875, Train loss: 0.5750768322944642 |Validation err: 0.318, Validation loss: 0.6768875233829021  
Epoch 7: Train err: 0.306, Train loss: 0.579937388420105 |Validation err: 0.3195, Validation loss: 0.6036276603117585  
Epoch 8: Train err: 0.302375, Train loss: 0.563445684671402 |Validation err: 0.346, Validation loss: 0.6068768594413996  
Epoch 9: Train err: 0.2955, Train loss: 0.5588727197647094 |Validation err: 0.3195, Validation loss: 0.601520698517561  
Epoch 10: Train err: 0.283125, Train loss: 0.545897435426712 |Validation err: 0.32, Validation loss: 0.5988965649157763  
Epoch 11: Train err: 0.275125, Train loss: 0.5378088033199311 |Validation err: 0.327, Validation loss: 0.619147676974535  
Epoch 12: Train err: 0.273625, Train loss: 0.535594067811966 |Validation err: 0.3535, Validation loss: 0.6480005066841841  
Epoch 13: Train err: 0.263625, Train loss: 0.5141938621997834 |Validation err: 0.3135, Validation loss: 0.6441182503476739  
Epoch 14: Train err: 0.249125, Train loss: 0.4992824411392212 |Validation err: 0.322, Validation loss: 0.6540486551821232  
Epoch 15: Train err: 0.249625, Train loss: 0.4969462950229645 |Validation err: 0.3275, Validation loss: 0.6473247641697526  
Epoch 16: Train err: 0.23325, Train loss: 0.4758114631175995 |Validation err: 0.3225, Validation loss: 0.6547067314386368  
Epoch 17: Train err: 0.257, Train loss: 0.5113468255996704 |Validation err: 0.3635, Validation loss: 0.6891361735761166  
Epoch 18: Train err: 0.264875, Train loss: 0.5206713597774506 |Validation err: 0.331, Validation loss: 0.6878945203498006  
Epoch 19: Train err: 0.234375, Train loss: 0.47404373025894164 |Validation err: 0.3405, Validation loss: 0.7093076854944229  
Epoch 20: Train err: 0.244625, Train loss: 0.4842326376438141 |Validation err: 0.3165, Validation loss: 0.7198585588485003  
Epoch 21: Train err: 0.234625, Train loss: 0.45815866231918334 |Validation err: 0.338, Validation loss: 0.712642149997377  
Epoch 22: Train err: 0.23125, Train loss: 0.45371274280548096 |Validation err: 0.3265, Validation loss: 0.749026388861239  
Epoch 23: Train err: 0.225875, Train loss: 0.464591454744339 |Validation err: 0.3405, Validation loss: 0.7113211695104837  
Epoch 24: Train err: 0.226625, Train loss: 0.44882566618919373 |Validation err: 0.331, Validation loss: 0.7977179307490587  
Epoch 25: Train err: 0.234375, Train loss: 0.4637635123729706 |Validation err: 0.3285, Validation loss: 0.7427194155752659  
Epoch 26: Train err: 0.240125, Train loss: 0.4794488083124161 |Validation err: 0.386, Validation loss: 0.7118616551160812  
Epoch 27: Train err: 0.250625, Train loss: 0.48315389394760133 |Validation err: 0.346, Validation loss: 0.8738224487751722  
Epoch 28: Train err: 0.23975, Train loss: 0.4596183142662048 |Validation err: 0.34, Validation loss: 0.8449452985078096  
Epoch 29: Train err: 0.218875, Train loss: 0.4428225445747376 |Validation err: 0.34, Validation loss: 0.8633094038814306  
Epoch 30: Train err: 0.233, Train loss: 0.46504191660881045 |Validation err: 0.352, Validation loss: 0.7883726954460144  
Finished Training  
Total time elapsed: 156.76 seconds



**Answer:**

It takes about 7 seconds longer than the default setting, but overall runtime remains similar. Raising the learning rate from 0.01 to 0.1 causes larger gradient-descent steps, so error and loss change more abruptly and overfitting appears earlier compared to the default model. As a result, the curves oscillate more and have sharper angles—indicating that a learning rate of 0.1 is too aggressive for this model.



### Part (c) - 3pt

Train `large_net` with all default parameters, including with `learning_rate=0.01`. Now, set `batch_size=512`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *increasing* the batch size.

```
In [20]: large_net = LargeNet()
train_net(large_net, 512, 0.01, 30)
large_path = get_model_name("large", batch_size=512, learning_rate=0.01, epoch=29)
plot_training_curve(large_path)
```



Epoch 1: Train err: 0.48175, Train loss: 0.6929379552602768 |Validation err: 0.478, Validation loss: 0.692682400345  
8023  
Epoch 2: Train err: 0.457625, Train loss: 0.692410409450531 |Validation err: 0.434, Validation loss: 0.691742524504  
6616  
Epoch 3: Train err: 0.437, Train loss: 0.6916500627994537 |Validation err: 0.4265, Validation loss: 0.6909130215644  
836  
Epoch 4: Train err: 0.433625, Train loss: 0.6908449903130531 |Validation err: 0.424, Validation loss: 0.68978703022  
00317  
Epoch 5: Train err: 0.434, Train loss: 0.6896935515105724 |Validation err: 0.424, Validation loss: 0.68813550472259  
52  
Epoch 6: Train err: 0.438, Train loss: 0.6883532106876373 |Validation err: 0.4285, Validation loss: 0.6860118657350  
54  
Epoch 7: Train err: 0.439375, Train loss: 0.6866871826350689 |Validation err: 0.426, Validation loss: 0.68369686603  
54614  
Epoch 8: Train err: 0.43525, Train loss: 0.6849770620465279 |Validation err: 0.411, Validation loss: 0.681467264890  
6708  
Epoch 9: Train err: 0.42375, Train loss: 0.6832008883357048 |Validation err: 0.414, Validation loss: 0.679591417312  
6221  
Epoch 10: Train err: 0.421, Train loss: 0.6811088025569916 |Validation err: 0.416, Validation loss: 0.6771541684865  
952  
Epoch 11: Train err: 0.42075, Train loss: 0.6794030703604221 |Validation err: 0.4095, Validation loss: 0.6748128235  
340118  
Epoch 12: Train err: 0.41475, Train loss: 0.6768062897026539 |Validation err: 0.412, Validation loss: 0.67370671033  
85925  
Epoch 13: Train err: 0.41025, Train loss: 0.6749698966741562 |Validation err: 0.412, Validation loss: 0.67061218619  
34662  
Epoch 14: Train err: 0.407125, Train loss: 0.6730909235775471 |Validation err: 0.4125, Validation loss: 0.669210046  
5297699  
Epoch 15: Train err: 0.4005, Train loss: 0.6706824786961079 |Validation err: 0.41, Validation loss: 0.6672518998384  
476  
Epoch 16: Train err: 0.397625, Train loss: 0.669177521020174 |Validation err: 0.405, Validation loss: 0.66490471363  
06763  
Epoch 17: Train err: 0.393875, Train loss: 0.6675704158842564 |Validation err: 0.401, Validation loss: 0.6630221307  
277679  
Epoch 18: Train err: 0.393, Train loss: 0.6648007929325104 |Validation err: 0.3945, Validation loss: 0.662391513586  
0443  
Epoch 19: Train err: 0.386125, Train loss: 0.6627416461706161 |Validation err: 0.3875, Validation loss: 0.659724175  
9300232  
Epoch 20: Train err: 0.38175, Train loss: 0.659615233540535 |Validation err: 0.401, Validation loss: 0.656416386365  
8905  
Epoch 21: Train err: 0.385875, Train loss: 0.6584861390292645 |Validation err: 0.388, Validation loss: 0.6586255580  
186844  
Epoch 22: Train err: 0.378625, Train loss: 0.6551117822527885 |Validation err: 0.3855, Validation loss: 0.652862340  
2118683  
Epoch 23: Train err: 0.371125, Train loss: 0.6508784741163254 |Validation err: 0.3835, Validation loss: 0.649810582  
3993683  
Epoch 24: Train err: 0.37675, Train loss: 0.6488191038370132 |Validation err: 0.384, Validation loss: 0.64748761057  
8537  
Epoch 25: Train err: 0.369, Train loss: 0.6445828042924404 |Validation err: 0.3825, Validation loss: 0.647354021668  
4341  
Epoch 26: Train err: 0.373, Train loss: 0.6428647711873055 |Validation err: 0.375, Validation loss: 0.6425962299108  
505  
Epoch 27: Train err: 0.35975, Train loss: 0.6372342072427273 |Validation err: 0.378, Validation loss: 0.63973273336  
88736  
Epoch 28: Train err: 0.354125, Train loss: 0.6337889917194843 |Validation err: 0.37, Validation loss: 0.64034420251  
84631  
Epoch 29: Train err: 0.353875, Train loss: 0.6311212778091431 |Validation err: 0.3665, Validation loss: 0.633566260  
3378296  
Epoch 30: Train err: 0.35275, Train loss: 0.6283531747758389 |Validation err: 0.367, Validation loss: 0.63244122266  
76941  
Finished Training  
Total time elapsed: 146.74 seconds



**Answer:**

With a larger batch size, training runs about 18 seconds faster since there are fewer iterations per epoch. Although the training error and loss are higher, the increased batch size reduces overfitting: the loss curves become smoother and the gap between training and validation curves narrows.


### Part (d) - 3pt

Train `large_net` with all default parameters, including with `learning_rate=0.01`. Now, set `batch_size=16`. Does the model take longer/shorter to train? Plot the training curve. Describe the effect of *decreasing* the batch size.

```
In [21]: large_net = LargeNet()
train_net(large_net, 16, 0.01, 30)

large_path = get_model_name("large", batch_size=16, learning_rate=0.01, epoch=29)
plot_training_curve(large_path)
```

Epoch 1: Train err: 0.432625, Train loss: 0.6775506126880646 |Validation err: 0.378, Validation loss: 0.6512571973800659  
Epoch 2: Train err: 0.366375, Train loss: 0.6387728816270828 |Validation err: 0.337, Validation loss: 0.612742235660553  
Epoch 3: Train err: 0.339375, Train loss: 0.6119522891640663 |Validation err: 0.345, Validation loss: 0.6396356213092804  
Epoch 4: Train err: 0.312125, Train loss: 0.5861616842746734 |Validation err: 0.3375, Validation loss: 0.6223122742176056  
Epoch 5: Train err: 0.309125, Train loss: 0.5655454085469246 |Validation err: 0.303, Validation loss: 0.5682719912528992  
Epoch 6: Train err: 0.283875, Train loss: 0.546434996843338 |Validation err: 0.3125, Validation loss: 0.581913836479187  
Epoch 7: Train err: 0.275625, Train loss: 0.5315411986708641 |Validation err: 0.3185, Validation loss: 0.5755203785896301  
Epoch 8: Train err: 0.253625, Train loss: 0.5110043309926987 |Validation err: 0.326, Validation loss: 0.6016001827716827  
Epoch 9: Train err: 0.249375, Train loss: 0.4988122125864029 |Validation err: 0.301, Validation loss: 0.5766945521831512  
Epoch 10: Train err: 0.234375, Train loss: 0.47769227081537247 |Validation err: 0.2935, Validation loss: 0.5778298568725586  
Epoch 11: Train err: 0.2245, Train loss: 0.4637914804518223 |Validation err: 0.295, Validation loss: 0.6023752751350403  
Epoch 12: Train err: 0.208625, Train loss: 0.4436814887523651 |Validation err: 0.3015, Validation loss: 0.5988463146686553  
Epoch 13: Train err: 0.1995, Train loss: 0.42190410897135733 |Validation err: 0.2965, Validation loss: 0.6023776261806488  
Epoch 14: Train err: 0.18225, Train loss: 0.3964607211649418 |Validation err: 0.285, Validation loss: 0.6563077869415284  
Epoch 15: Train err: 0.176625, Train loss: 0.38504845155775547 |Validation err: 0.3015, Validation loss: 0.684573037981987  
Epoch 16: Train err: 0.160125, Train loss: 0.35515162971615793 |Validation err: 0.313, Validation loss: 0.7153961782455445  
Epoch 17: Train err: 0.146875, Train loss: 0.33127270932495595 |Validation err: 0.3095, Validation loss: 0.8292664550542831  
Epoch 18: Train err: 0.1435, Train loss: 0.3216609486192465 |Validation err: 0.3265, Validation loss: 0.8005202133655548  
Epoch 19: Train err: 0.130375, Train loss: 0.30764051334559916 |Validation err: 0.3255, Validation loss: 0.8987283756732941  
Epoch 20: Train err: 0.11675, Train loss: 0.28437252435833216 |Validation err: 0.309, Validation loss: 0.8813227322101593  
Epoch 21: Train err: 0.121625, Train loss: 0.27901479678601027 |Validation err: 0.3175, Validation loss: 0.8441186798810959  
Epoch 22: Train err: 0.112875, Train loss: 0.2728129132091999 |Validation err: 0.3125, Validation loss: 0.8382637614011764  
Epoch 23: Train err: 0.1, Train loss: 0.24255631332099437 |Validation err: 0.3265, Validation loss: 1.0690370055437088  
Epoch 24: Train err: 0.095375, Train loss: 0.2318725396282971 |Validation err: 0.322, Validation loss: 1.0889509381055833  
Epoch 25: Train err: 0.095875, Train loss: 0.23257503168098628 |Validation err: 0.329, Validation loss: 1.0898759965896607  
Epoch 26: Train err: 0.08275, Train loss: 0.20350570978596808 |Validation err: 0.321, Validation loss: 1.127944498538971  
Epoch 27: Train err: 0.09075, Train loss: 0.2147584371343255 |Validation err: 0.3205, Validation loss: 1.1985667142868042  
Epoch 28: Train err: 0.076125, Train loss: 0.18171320636058227 |Validation err: 0.322, Validation loss: 1.3782274899482727  
Epoch 29: Train err: 0.091875, Train loss: 0.21968170262500644 |Validation err: 0.3245, Validation loss: 1.2601391798257828  
Epoch 30: Train err: 0.07725, Train loss: 0.1865569370612502 |Validation err: 0.323, Validation loss: 1.3147952773571014  
Finished Training  
Total time elapsed: 225.82 seconds



Answer:

This takes around 61 seconds longer, decreasing the batch size increases training time since more iterations are needed per epoch. It does, however, lower the training error and loss and the overfitting occurs earlier. For validation, the error and loss is still large.

## Part 4. Hyperparameter Search [6 pt]

### Part (a) - 2pt

Based on the plots from above, choose another set of values for the hyperparameters (network, batch\_size, learning\_rate) that you think would help you improve the validation accuracy. Justify your choice.

**Answer:**

I would chose large network (with more parameters), large batch size, and small learning rate. So, I choose network = large\_net; batch\_size = 128, learning\_rate = 0.01, and epoch = 30.


The large network outperforms the small network thanks to its greater parameter capacity. From Part 3, a learning rate of 0.01 proved to be well balanced. Small batch sizes tend to cause overfitting, so using a larger batch improves performance without overfitting. Finally, training for 30 epochs strikes a good balance between underfitting and overfitting.

## Part (b) - 1pt

Train the model with the hyperparameters you chose in part(a), and include the training curve.

```
In [22]: large_net = LargeNet()
train_net(large_net, 128, 0.01, 30)
model_path_large = get_model_name("large", batch_size=128, learning_rate=0.01, epoch=29)
plot_training_curve(model_path_large)
```

Epoch 1: Train err: 0.454375, Train loss: 0.6920222280517457 |Validation err: 0.426, Validation loss: 0.6897575519979  
Epoch 2: Train err: 0.44475, Train loss: 0.6877915528085496 |Validation err: 0.411, Validation loss: 0.6820644401013851  
Epoch 3: Train err: 0.422125, Train loss: 0.6812269706574697 |Validation err: 0.412, Validation loss: 0.6738929860293865  
Epoch 4: Train err: 0.41025, Train loss: 0.673797849624876 |Validation err: 0.4075, Validation loss: 0.6660339944064617  
Epoch 5: Train err: 0.397375, Train loss: 0.6661591208170331 |Validation err: 0.4015, Validation loss: 0.6590236835181713  
Epoch 6: Train err: 0.377125, Train loss: 0.6567326358386448 |Validation err: 0.3895, Validation loss: 0.6513406559824944  
Epoch 7: Train err: 0.374625, Train loss: 0.6487630000190129 |Validation err: 0.3765, Validation loss: 0.639539323747158  
Epoch 8: Train err: 0.36025, Train loss: 0.6365784785104176 |Validation err: 0.372, Validation loss: 0.6431583352386951  
Epoch 9: Train err: 0.35, Train loss: 0.6283171810801067 |Validation err: 0.346, Validation loss: 0.6241483055055141  
Epoch 10: Train err: 0.34875, Train loss: 0.6217929465430123 |Validation err: 0.346, Validation loss: 0.6209967844188213  
Epoch 11: Train err: 0.333, Train loss: 0.6105351911650764 |Validation err: 0.341, Validation loss: 0.6170116998255253  
Epoch 12: Train err: 0.327125, Train loss: 0.5992055573160686 |Validation err: 0.327, Validation loss: 0.6038246862590313  
Epoch 13: Train err: 0.32225, Train loss: 0.5948410185556563 |Validation err: 0.319, Validation loss: 0.5991976037621498  
Epoch 14: Train err: 0.307875, Train loss: 0.5794498173017351 |Validation err: 0.33, Validation loss: 0.614487636834383  
Epoch 15: Train err: 0.304125, Train loss: 0.5702668012134613 |Validation err: 0.3365, Validation loss: 0.6122750714421272  
Epoch 16: Train err: 0.295375, Train loss: 0.5657135503632682 |Validation err: 0.3295, Validation loss: 0.6064109429717064  
Epoch 17: Train err: 0.29525, Train loss: 0.5594405839367519 |Validation err: 0.3125, Validation loss: 0.5987546071410179  
Epoch 18: Train err: 0.28475, Train loss: 0.5496182214646113 |Validation err: 0.31, Validation loss: 0.5874427258968353  
Epoch 19: Train err: 0.275, Train loss: 0.5384441462774125 |Validation err: 0.323, Validation loss: 0.5953904837369919  
Epoch 20: Train err: 0.276, Train loss: 0.5352702755776663 |Validation err: 0.3065, Validation loss: 0.593407429754734  
Epoch 21: Train err: 0.2675, Train loss: 0.5263815578960237 |Validation err: 0.3025, Validation loss: 0.5831845328211784  
Epoch 22: Train err: 0.263375, Train loss: 0.5202986276338971 |Validation err: 0.308, Validation loss: 0.5872550122439861  
Epoch 23: Train err: 0.257875, Train loss: 0.5148504934613667 |Validation err: 0.3055, Validation loss: 0.5826834626495838  
Epoch 24: Train err: 0.257125, Train loss: 0.5102080371644762 |Validation err: 0.3165, Validation loss: 0.6045809499919415  
Epoch 25: Train err: 0.249, Train loss: 0.5024650309767041 |Validation err: 0.3035, Validation loss: 0.5960421841591597  
Epoch 26: Train err: 0.244, Train loss: 0.491200448028625 |Validation err: 0.3065, Validation loss: 0.5820110738277435  
Epoch 27: Train err: 0.243875, Train loss: 0.491646780381127 |Validation err: 0.3045, Validation loss: 0.5974767580628395  
Epoch 28: Train err: 0.23975, Train loss: 0.4839872277918316 |Validation err: 0.3015, Validation loss: 0.5878518670797348  
Epoch 29: Train err: 0.2285, Train loss: 0.4747762259036776 |Validation err: 0.3215, Validation loss: 0.6145458891987801  
Epoch 30: Train err: 0.227375, Train loss: 0.4695242298027826 |Validation err: 0.307, Validation loss: 0.5974648240953684  
Finished Training  
Total time elapsed: 152.66 seconds



### Part (c) - 2pt

Based on your result from Part(a), suggest another set of hyperparameter values to try. Justify your choice.

**Answer:**

I'll use the large network with a batch size of 128, a learning rate of 0.005, and train for 50 epochs. I increased the epoch count to 50 to get more accurate data with less validation loss and error due to the smaller learning rate. Lowering the learning rate allows the model to make more precise parameter updates.

### Part (d) - 1pt

Train the model with the hyperparameters you chose in part(c), and include the training curve.

```
In [26]: large_net = LargeNet()
train_net(large_net, 128, 0.005, 50)
model_path_large = get_model_name("large", batch_size=128, learning_rate=0.005, epoch=49)
plot_training_curve(model_path_large)
```

Epoch 1: Train err: 0.49775, Train loss: 0.6930236144671365 |Validation err: 0.5095, Validation loss: 0.69302840903  
40137  
Epoch 2: Train err: 0.45725, Train loss: 0.6922709601266044 |Validation err: 0.42, Validation loss: 0.6916764788329  
601  
Epoch 3: Train err: 0.433375, Train loss: 0.6913018368539356 |Validation err: 0.4115, Validation loss: 0.6902153231  
203556  
Epoch 4: Train err: 0.427875, Train loss: 0.6897292752114553 |Validation err: 0.418, Validation loss: 0.68726608529  
68693  
Epoch 5: Train err: 0.4315, Train loss: 0.6873460137654864 |Validation err: 0.418, Validation loss: 0.6837649568915  
367  
Epoch 6: Train err: 0.427625, Train loss: 0.6840077771080865 |Validation err: 0.415, Validation loss: 0.67932287231  
08768  
Epoch 7: Train err: 0.4185, Train loss: 0.679767336164202 |Validation err: 0.4035, Validation loss: 0.6747710406780  
243  
Epoch 8: Train err: 0.40525, Train loss: 0.6748903526200188 |Validation err: 0.398, Validation loss: 0.671427298337  
2211  
Epoch 9: Train err: 0.398875, Train loss: 0.6694405656012278 |Validation err: 0.3975, Validation loss: 0.6635362096  
130848  
Epoch 10: Train err: 0.38825, Train loss: 0.6631133367144872 |Validation err: 0.395, Validation loss: 0.65923252701  
75934  
Epoch 11: Train err: 0.379875, Train loss: 0.6555238801335531 |Validation err: 0.3895, Validation loss: 0.654037024  
8258114  
Epoch 12: Train err: 0.377125, Train loss: 0.6498739918073019 |Validation err: 0.3755, Validation loss: 0.648442260  
9210014  
Epoch 13: Train err: 0.373625, Train loss: 0.6444921550296602 |Validation err: 0.39, Validation loss: 0.65095221996  
30737  
Epoch 14: Train err: 0.365375, Train loss: 0.6393705587538462 |Validation err: 0.3825, Validation loss: 0.647908397  
0189095  
Epoch 15: Train err: 0.36475, Train loss: 0.6360937640780494 |Validation err: 0.38, Validation loss: 0.645784851163  
6257  
Epoch 16: Train err: 0.358875, Train loss: 0.6319135682923454 |Validation err: 0.3695, Validation loss: 0.641961343  
5864449  
Epoch 17: Train err: 0.35625, Train loss: 0.6300226449966431 |Validation err: 0.3855, Validation loss: 0.6470319963  
991642  
Epoch 18: Train err: 0.348375, Train loss: 0.625270790523953 |Validation err: 0.3705, Validation loss: 0.6400919146  
835804  
Epoch 19: Train err: 0.346625, Train loss: 0.6228595243559943 |Validation err: 0.364, Validation loss: 0.6333517804  
741859  
Epoch 20: Train err: 0.340375, Train loss: 0.6144135934965951 |Validation err: 0.356, Validation loss: 0.6322385892  
271996  
Epoch 21: Train err: 0.335, Train loss: 0.6116347029095605 |Validation err: 0.355, Validation loss: 0.6274648457765  
579  
Epoch 22: Train err: 0.327125, Train loss: 0.6046467820803324 |Validation err: 0.355, Validation loss: 0.6266518495  
976925  
Epoch 23: Train err: 0.323125, Train loss: 0.5985852300174652 |Validation err: 0.344, Validation loss: 0.6204004175  
961018  
Epoch 24: Train err: 0.326125, Train loss: 0.5954039200903878 |Validation err: 0.3335, Validation loss: 0.617544036  
3585949  
Epoch 25: Train err: 0.311625, Train loss: 0.5851705992032611 |Validation err: 0.3285, Validation loss: 0.611986499  
2797375  
Epoch 26: Train err: 0.3065, Train loss: 0.5780152792022342 |Validation err: 0.3265, Validation loss: 0.60789039358  
49667  
Epoch 27: Train err: 0.304875, Train loss: 0.5721108657973153 |Validation err: 0.3255, Validation loss: 0.608403999  
3584156  
Epoch 28: Train err: 0.29775, Train loss: 0.5698803469302163 |Validation err: 0.3355, Validation loss: 0.6119430214  
166641  
Epoch 29: Train err: 0.2955, Train loss: 0.5603470627277617 |Validation err: 0.3235, Validation loss: 0.62185088917  
61303  
Epoch 30: Train err: 0.28825, Train loss: 0.5553928585279555 |Validation err: 0.3205, Validation loss: 0.6035062558  
948994  
Epoch 31: Train err: 0.286625, Train loss: 0.5519701001190004 |Validation err: 0.342, Validation loss: 0.6221818365  
15665  
Epoch 32: Train err: 0.276875, Train loss: 0.5457092022138929 |Validation err: 0.3175, Validation loss: 0.588867370  
0392246  
Epoch 33: Train err: 0.283125, Train loss: 0.5443956478247567 |Validation err: 0.317, Validation loss: 0.5931047797  
203064  
Epoch 34: Train err: 0.272875, Train loss: 0.533677642780637 |Validation err: 0.306, Validation loss: 0.59393728524  
44649  
Epoch 35: Train err: 0.26825, Train loss: 0.5300850234334431 |Validation err: 0.3065, Validation loss: 0.5930652040  
988207  
Epoch 36: Train err: 0.264125, Train loss: 0.5256856509617397 |Validation err: 0.304, Validation loss: 0.5851593222  
469091  
Epoch 37: Train err: 0.261625, Train loss: 0.5227410575700184 |Validation err: 0.308, Validation loss: 0.5821809023  
618698  
Epoch 38: Train err: 0.262375, Train loss: 0.5190259246599107 |Validation err: 0.301, Validation loss: 0.5799242295  
324802  
Epoch 39: Train err: 0.2565, Train loss: 0.5132552896227155 |Validation err: 0.312, Validation loss: 0.602854173630

476  
Epoch 40: Train err: 0.254125, Train loss: 0.5053303662746672 |Validation err: 0.3085, Validation loss: 0.5967046171426773  
Epoch 41: Train err: 0.255, Train loss: 0.5049237902202304 |Validation err: 0.2945, Validation loss: 0.5929372347891331  
Epoch 42: Train err: 0.251, Train loss: 0.5055604023592812 |Validation err: 0.295, Validation loss: 0.5962063446640968  
Epoch 43: Train err: 0.249625, Train loss: 0.4992753877526238 |Validation err: 0.3015, Validation loss: 0.5894495025277138  
Epoch 44: Train err: 0.250125, Train loss: 0.5041983591185676 |Validation err: 0.33, Validation loss: 0.6286768466234207  
Epoch 45: Train err: 0.24425, Train loss: 0.4957300717868502 |Validation err: 0.309, Validation loss: 0.5885458998382092  
Epoch 46: Train err: 0.237625, Train loss: 0.4858499401145511 |Validation err: 0.3105, Validation loss: 0.5847885571420193  
Epoch 47: Train err: 0.243625, Train loss: 0.4965810960247403 |Validation err: 0.2955, Validation loss: 0.5891992636024952  
Epoch 48: Train err: 0.233, Train loss: 0.47616930518831524 |Validation err: 0.299, Validation loss: 0.5792636536061764  
Epoch 49: Train err: 0.2345, Train loss: 0.4818397746199653 |Validation err: 0.296, Validation loss: 0.5824931841343641  
Epoch 50: Train err: 0.229, Train loss: 0.47108897897932267 |Validation err: 0.303, Validation loss: 0.5918887332081795  
Finished Training  
Total time elapsed: 240.19 seconds







## Part 5. Evaluating the Best Model [15 pt]

### Part (a) - 1pt

Choose the **best** model that you have so far. This means choosing the best model checkpoint, including the choice of `small_net` vs `large_net`, the `batch_size`, `learning_rate`, and the **epoch number**.

Modify the code below to load your chosen set of weights to the model object `net`.

```
In [27]: net = LargeNet()
model_path = get_model_name(net.name, batch_size=128, learning_rate=0.005, epoch=49)
state = torch.load(model_path)
net.load_state_dict(state)
```

```
Out[27]: <All keys matched successfully>
```

### Part (b) - 2pt

Justify your choice of model from part (a).

**Answer:**

- Validation error from part(a): 0.303
- Validation loss from part(a): 0.591

Both the validation loss and error here are better than previous results.

I'm selecting the large network because its higher capacity yields lower validation error than the small network. Setting the batch size to 128 and the learning rate to 0.005 significantly reduces overfitting. With this lower learning rate, extending training to 50 epochs further refines the model, minimizing validation loss and error.

### Part (c) - 2pt

Using the code in Part 0, any code from lecture notes, or any code that you write, compute and report the **test classification error** for your chosen model.

```
In [32]: # If you use the `evaluate` function provided in part 0, you will need to
# set batch_size > 1
train_loader, val_loader, test_loader, classes = get_data_loader(
    target_classes=["cat", "dog"],
    batch_size=128)
```

```
criterion = nn.BCEWithLogitsLoss()
test_err, test_loss = evaluate(net, test_loader, criterion)
print("The test classification error and loss are :", test_err, "and", test_loss)
```

The test classification error and loss are : 0.2855 and 0.5665808189660311

## Part (d) - 3pt

How does the test classification error compare with the **validation error**? Explain why you would expect the test error to be *higher* than the validation error.

**Answer:**

The test error (0.2855) is actually slightly lower than the validation error (0.303). In practice, one usually expects the test error to be a bit higher because hyperparameters are tuned to minimize validation loss—so the model has “seen” the validation set indirectly via that tuning. Here, though, the difference is small and likely just sampling noise: the particular split used for testing happened to be easier than the validation split.

## Part (e) - 2pt

Why did we only use the test data set at the very end? Why is it important that we use the test data as little as possible?

**Answer:**

We reserve the test set for the very end because it has never been used for training or validation, making it a true simulation of real-world unseen data and an unbiased measure of final model performance. If we evaluate or tune on the test set during development, it effectively becomes part of the training process—leading to overfitting, where the model “remembers” the test examples rather than generalizing to new data.

## Part (f) - 5pt

How does the your best CNN model compare with an 2-layer ANN model (no convolutional layers) on classifying cat and dog images. You can use a 2-layer ANN architecture similar to what you used in Lab 1. You should explore different hyperparameter settings to determine how well you can do on the validation dataset. Once satisfied with the performance, you may test it out on the test data.

Hint: The ANN in lab 1 was applied on greyscale images. The cat and dog images are colour (RGB) and so you will need to flatted and concatinete all three colour layers before feeding them into an ANN.

```
In [33]: torch.manual_seed(1) #set random seed
class CatAndDog(nn.Module): # 2-layer artificial neural network
    def __init__(self):
        self.name = "CatAndDog"
        super(CatAndDog, self).__init__()
        self.layer1 = nn.Linear(3 * 32 * 32, 60)
        self.layer2 = nn.Linear(60, 1)
    def forward(self, img):
        flattened = img.view(-1, 3 * 32 * 32)
        activation1 = self.layer1(flattened)
        activation1 = F.relu(activation1)
        activation2 = self.layer2(activation1)
        activation2 = activation2.squeeze(1)
        return activation2

train_ANN = CatAndDog()
train_net(train_ANN, batch_size=128, learning_rate=0.005, num_epochs=50)


cdpath = get_model_name("CatAndDog", batch_size=128, learning_rate=0.005, epoch=49)
plot_training_curve(cdpath)
```

Epoch 1: Train err: 0.420125, Train loss: 0.673239356941647 |Validation err: 0.405, Validation loss: 0.661061722785  
2345  
Epoch 2: Train err: 0.382625, Train loss: 0.651417494766296 |Validation err: 0.397, Validation loss: 0.653374370187  
521  
Epoch 3: Train err: 0.369125, Train loss: 0.6399649864151364 |Validation err: 0.3865, Validation loss: 0.6498257815  
83786  
Epoch 4: Train err: 0.36025, Train loss: 0.6315847427125961 |Validation err: 0.396, Validation loss: 0.653963312506  
6757  
Epoch 5: Train err: 0.35025, Train loss: 0.6255956140775529 |Validation err: 0.3855, Validation loss: 0.64519455283  
88023  
Epoch 6: Train err: 0.3365, Train loss: 0.6164521328986637 |Validation err: 0.3805, Validation loss: 0.648137576878  
0708  
Epoch 7: Train err: 0.329625, Train loss: 0.6098158993418255 |Validation err: 0.3795, Validation loss: 0.6461284793  
913364  
Epoch 8: Train err: 0.32, Train loss: 0.6008976754688081 |Validation err: 0.368, Validation loss: 0.642281174659729  
Epoch 9: Train err: 0.31875, Train loss: 0.5951670493398394 |Validation err: 0.3785, Validation loss: 0.64245909824  
96738  
Epoch 10: Train err: 0.30975, Train loss: 0.587331097277384 |Validation err: 0.3685, Validation loss: 0.64607052877  
54536  
Epoch 11: Train err: 0.300625, Train loss: 0.5775917257581439 |Validation err: 0.373, Validation loss: 0.6475792638  
9575  
Epoch 12: Train err: 0.295, Train loss: 0.5687684615453085 |Validation err: 0.3565, Validation loss: 0.643708620220  
4227  
Epoch 13: Train err: 0.292125, Train loss: 0.5605954963063436 |Validation err: 0.3795, Validation loss: 0.653732180  
595398  
Epoch 14: Train err: 0.2805, Train loss: 0.5532135897212558 |Validation err: 0.3775, Validation loss: 0.65532146021  
72375  
Epoch 15: Train err: 0.26975, Train loss: 0.5424539191382272 |Validation err: 0.365, Validation loss: 0.65238736569  
88144  
Epoch 16: Train err: 0.267625, Train loss: 0.5364074565115429 |Validation err: 0.3575, Validation loss: 0.644939459  
8603249  
Epoch 17: Train err: 0.257875, Train loss: 0.5254545476701524 |Validation err: 0.3625, Validation loss: 0.646231323  
4806061  
Epoch 18: Train err: 0.255625, Train loss: 0.5225144790278541 |Validation err: 0.3615, Validation loss: 0.655661832  
5412273  
Epoch 19: Train err: 0.24275, Train loss: 0.5108178385666439 |Validation err: 0.3685, Validation loss: 0.6728108972  
31102  
Epoch 20: Train err: 0.2415, Train loss: 0.49850498636563617 |Validation err: 0.3635, Validation loss: 0.6538236998  
021603  
Epoch 21: Train err: 0.224625, Train loss: 0.48706356353229946 |Validation err: 0.3705, Validation loss: 0.65819843  
11342239  
Epoch 22: Train err: 0.224875, Train loss: 0.4807430718626295 |Validation err: 0.373, Validation loss: 0.6835006698  
966026  
Epoch 23: Train err: 0.215875, Train loss: 0.46818161625710747 |Validation err: 0.3695, Validation loss: 0.65887328  
60982418  
Epoch 24: Train err: 0.198625, Train loss: 0.45491003138678415 |Validation err: 0.352, Validation loss: 0.663257122  
0397949  
Epoch 25: Train err: 0.196, Train loss: 0.4448101203592997 |Validation err: 0.3665, Validation loss: 0.671755049377  
6798  
Epoch 26: Train err: 0.192625, Train loss: 0.4344807347607991 |Validation err: 0.364, Validation loss: 0.6809879615  
902901  
Epoch 27: Train err: 0.18825, Train loss: 0.42924191270555767 |Validation err: 0.36, Validation loss: 0.70343190431  
59485  
Epoch 28: Train err: 0.186125, Train loss: 0.41850896487160333 |Validation err: 0.352, Validation loss: 0.696032881  
7367554  
Epoch 29: Train err: 0.174, Train loss: 0.4000347199894133 |Validation err: 0.362, Validation loss: 0.6942745745182  
037  
Epoch 30: Train err: 0.159375, Train loss: 0.38698555363549125 |Validation err: 0.3525, Validation loss: 0.69802417  
23358631  
Epoch 31: Train err: 0.151625, Train loss: 0.3733342725133139 |Validation err: 0.3505, Validation loss: 0.693617589  
7717476  
Epoch 32: Train err: 0.156625, Train loss: 0.374767557969169 |Validation err: 0.362, Validation loss: 0.73524991422  
89162  
Epoch 33: Train err: 0.1545, Train loss: 0.3645764111526429 |Validation err: 0.351, Validation loss: 0.739649277180  
4333  
Epoch 34: Train err: 0.14325, Train loss: 0.3538515600893233 |Validation err: 0.358, Validation loss: 0.71553310379  
38595  
Epoch 35: Train err: 0.12675, Train loss: 0.329066820087887 |Validation err: 0.3555, Validation loss: 0.72628996521  
23451  
Epoch 36: Train err: 0.132875, Train loss: 0.3315145470320232 |Validation err: 0.363, Validation loss: 0.7645671069  
62204  
Epoch 37: Train err: 0.121375, Train loss: 0.3176128679797763 |Validation err: 0.3545, Validation loss: 0.742336049  
6759415  
Epoch 38: Train err: 0.10825, Train loss: 0.2981732565732229 |Validation err: 0.355, Validation loss: 0.75149273499  
84646  
Epoch 39: Train err: 0.106125, Train loss: 0.29321087636644877 |Validation err: 0.36, Validation loss: 0.7950831502  
67601

Epoch 40: Train err: 0.1125, Train loss: 0.29049530294206405 |Validation err: 0.355, Validation loss: 0.7926928251981735  
Epoch 41: Train err: 0.08975, Train loss: 0.2697687998177513 |Validation err: 0.3505, Validation loss: 0.786663856357336  
Epoch 42: Train err: 0.091875, Train loss: 0.2634454866250356 |Validation err: 0.368, Validation loss: 0.8070601522922516  
Epoch 43: Train err: 0.084375, Train loss: 0.24909782291404783 |Validation err: 0.3585, Validation loss: 0.8075341656804085  
Epoch 44: Train err: 0.08275, Train loss: 0.2445363224971862 |Validation err: 0.362, Validation loss: 0.8047830387949944  
Epoch 45: Train err: 0.074375, Train loss: 0.2346846525158201 |Validation err: 0.353, Validation loss: 0.8283248022198677  
Epoch 46: Train err: 0.071375, Train loss: 0.22386179202132755 |Validation err: 0.376, Validation loss: 0.8855374231934547  
Epoch 47: Train err: 0.068875, Train loss: 0.21881483873677632 |Validation err: 0.3565, Validation loss: 0.8784874901175499  
Epoch 48: Train err: 0.0695, Train loss: 0.2134316246660929 |Validation err: 0.3555, Validation loss: 0.8468413315713406  
Epoch 49: Train err: 0.06, Train loss: 0.19554019305441114 |Validation err: 0.3665, Validation loss: 0.8748635873198509  
Epoch 50: Train err: 0.057625, Train loss: 0.1915858000438388 |Validation err: 0.358, Validation loss: 0.9007941670715809  
Finished Training  
Total time elapsed: 215.55 seconds



Train vs Validation Loss



```
In [34]: train_loader, val_loader, test_loader, classes = get_data_loader(
        target_classes=["cat", "dog"],
        batch_size=128)

error, loss = evaluate(train_ANN, test_loader, nn.BCEWithLogitsLoss())
print("Test error is: ", error, ", and test loss is: ", loss)

error1, loss1 = evaluate(train_ANN, val_loader, nn.BCEWithLogitsLoss())
print("Validation error: ", error1, ", and validation loss: ", loss1)
```

Test error is: 0.3565 , and test loss is: 0.9038791991770267  
 Validation error: 0.358 , and validation loss: 0.9007261469960213

#### Answer:

The ANN model shows higher test error and loss, so the CNN model performs better.